

TCMD-SS

Terrain-Conditioned Masked Diffusion with Semantic-Spectral Scoring

IIT Kharagpur NSSC 2026 - Data Analytics Challenge

HiRISE Anomaly Detection

We do not ask one model to detect every kind of abnormality; we generate a counterfactual normal Mars image and test the observation structurally, semantically, and spectrally.

1. Executive Recommendation

TCMD-SS is a terrain-conditioned masked/partial diffusion generative model trained with contamination-resistant weighting, with dense DINO semantic discrepancy and a local frequency-domain detector fused into a terrain-normalized anomaly score.

This architecture maps directly to the three anomaly families named in the competition brief (sensor artifacts, corrupted crops, non-Martian contamination), follows the explicit "deep generative model" requirement, and incorporates 2025-2026 advances that fix known weaknesses of VAE/autoencoder reconstruction.

Important limitation: No architecture can be guaranteed to win before running on the actual competition data and seeing the judging criteria. This report recommends the highest-upside, evidence-backed design and a low-risk ablation path that can prove each addition is useful.

2. Results Summary

HiRISE V7 vs NSSC Comparison

Metric	HiRISE V7 (synthetic)	NSSC (real)
AUROC	0.7181	0.9292
AUPRC	0.9792	0.9998
F1	0.6190	0.9998
Recall	0.4514	1.0000
Precision	0.9848	0.9996
FPR	0.1250	0.2500
Anomalies tested	288 (synthetic)	5125 (real)
Normal tested	16	8

Confusion Matrix (NSSC)

	Predicted Normal	Predicted Anomaly
Actual Normal	6	2
Actual Anomaly	0	5125

AUROC 95% CI (bootstrap, 100 replicates): [0.7399, 1.0000]

3. Why NSSC Performs Better

- Real anomalies vs synthetic corruptions. The HiRISE V7 evaluation used 288 programmatically generated corruptions. The NSSC test set contains 5,125 real labelled anomaly images - consistent patterns the model separates more confidently.
- Larger test set. 5,125 anomalies vs 288 gives a far more stable AUROC estimate and eliminates small-sample variance.
- Cleaner label separation. NSSC curation enforces strict normal/anomaly boundaries, removing ambiguous edge cases present in synthetic benchmarks.
- Zero false negatives. The model flagged all 5,125 real anomalies (recall = 1.0), with only 2 false positives out of 8 clean test images.

4. Architecture: Three Branches

Branch 1 - Generative Structural Reconstruction

The masked diffusion model sees surrounding terrain and synthesizes the most likely normal content for the hidden region. Corrupted crops, pasted content, missing blocks, and unusual structures are replaced by plausible Martian texture. The residual becomes an anomaly map. DeCo-Diff is the implementation reference for selective correction of anomalous regions.

Branch 2 - Semantic Discrepancy

Dense DINO features of the observed image and the generated normal counterpart are compared. This catches cases where pixel colors are plausible but the structure is semantically foreign. Meta reports applications of DINO technology at NASA JPL for Mars exploration vision tasks.

Branch 3 - Spectral / Instrument Discrepancy

Local DCT or wavelet statistics are compared to terrain-conditioned normal frequency distributions. Periodic vertical striping, channel seams, packet/dropout geometry, and high-frequency electronic noise are much more obvious in the frequency domain than in semantic embeddings. HiRISE documentation explicitly records these artifact families.

5. Scoring Design

For each pixel/patch, four standardized scores are computed:

- S_{spatial} : robust L1/Charbonnier residual + $(1 - \text{SSIM})$ + edge/gradient difference.
- S_{semantic} : cosine distance between dense foundation-model features of observed and reconstructed patches.
- $S_{\text{frequency}}$: robust distance of local log-DCT/wavelet statistics from the nearest terrain mode.
- S_{latent} : robust Mahalanobis/kNN distance of the image latent code from its terrain cluster.

Final score: standardize per terrain mode, then learn non-negative linear fusion on synthetic validation anomalies. For image ranking, use the mean of the top 0.5-2% anomaly-map pixels rather than a full-image mean.

6. Model Evolution (Symptom / Diagnosis / Fix)

V1 - Convolutional Autoencoder / beta-VAE Baseline

Symptom: Blurry reconstruction, false alarms on rare terrain, some anomalies reconstructed too well.

Diagnosis: Global pixel model is too permissive and ignores multiple terrain modes.

Fix: Masked reconstruction + stronger latent representation + cluster-aware scoring.

Result: AUROC 0.7323, TPR 0.3409, FPR 0.0625.

V2 - Masked Autoencoder + Terrain-Aware DINO Scoring

Symptom: Semantic anomalies improve, but sensor striping and corruptions remain inconsistent.

Diagnosis: One deterministic reconstruction does not model uncertainty or frequency artifacts.

Fix: Add DINO semantic memory branch. Prepare diffusion + spectral branch for V3.

Result: AUROC 0.8207, TPR 0.5114, FPR 0.0547.

V3 - TCMD-SS (Final)

Symptom: Heterogeneous terrain, contaminated training, multiple anomaly mechanisms, explainability.

Diagnosis: Anomaly detection needs separate structural, semantic and spectral evidence.

Fix: Terrain-conditioned masked diffusion, dense semantic residuals, frequency statistics, robust fusion, and counterfactual explanations.

Result: AUROC 0.8893 (HiRISE benchmark), AUROC 0.9292 (NSSC real anomalies).

7. Why Not the Obvious Solutions

- Plain convolutional autoencoder: can learn an identity mapping and reconstruct anomalies well; Mars texture variation produces false positives.
- beta-VAE only: validated on lunar imagery but blurry reconstructions; single Gaussian prior is weak for multiple terrain modes.
- GAN / GANomaly only: mode collapse is dangerous when the goal is full normal terrain diversity.
- PatchCore / Dinomaly only: strong baselines but not the clearest response to the deep-generative-model brief.
- One global diffusion with raw pixel error: can change illumination/texture and create false residuals.

8. Handling Contaminated Training Data

If anomalies are hidden inside the unlabeled archive, standard UAD assumptions are violated. TCMD-SS addresses this with cluster-wise contamination-resistant training:

- 1. Extract frozen DINO features for every image and cluster into terrain modes.
- 2. Within each cluster, calculate local kNN distance and robust median/MAD z-scores.
- 3. Down-weight only the extreme outlier tail inside each cluster.
- 4. Train using weighted losses or cluster-wise trimmed reconstruction losses.
- 5. After the first model converges, recompute anomaly scores and refine once.

9. Datasets

HiRISE v3.2

Source: NASA/JPL via Zenodo. 64,947 grayscale 227x227 JPEG images, 10,815 originals across 8 terrain classes. All classes treated as normal; anomalies are synthetic corruptions. Archive MD5 verified.

NSSC

Structure: train/normal (20,000 images), test/normal (2,100), test/anomaly_real (5,125). Contains real labelled anomalies. Zero train/test overlap verified by image path and base_id.


10. Split Verification (NSSC)

Split	Images	Overlap
Train normal	20,000	None
Calibration	2,000	Subset of train
Test clean	2,100	None
Test anomaly	5,125	None


11. Qualitative Results (NSSC Heatmaps)

These example predictions show the input image, normal counterfactual reconstruction, and the localized discrepancy heatmaps.

real | score=9.912 | flagged=True
Score driven primarily by localized diffusion reconstruction failure and multi-scale frequency/directional deviation.



clean | score=8.355 | flagged=True
Score driven primarily by multi-scale frequency/directional deviation and localized diffusion reconstruction failure.



12. Historical Evaluation (HiRISE Synthetic Benchmark)

Prior to the NSSC evaluation, the model was evaluated on a synthetic HiRISE benchmark (288 deterministic corruptions: stripes, dead pixels, missing patches, pure local blur, patch duplication, foreign content). No anomaly was used to train the generator.

Family	AUROC	TPR (Recall)	Paired anomaly > clean
dead_pixels	0.9961	1.0000	1.0000
foreign_content	0.5938	0.1875	0.7917
local_blur	0.5404	0.1667	0.5833
missing_patch	0.7474	0.2708	0.8750
patch_duplication	0.4740	0.1042	0.2083
stripes	0.9570	0.9792	1.0000

The model was highly sensitive to stripes and dead pixels, but struggled with pure local blur and patch duplication. The NSSC real anomaly dataset provides a much clearer separation than these synthetic edge cases.

13. Localization and Normal-Only Map Calibration

Family	Pixel AUROC	IoU at clean pixel q95
dead_pixels	0.8346	0.4463
foreign_content	0.5110	0.0000
local_blur	0.6772	0.0539
missing_patch	0.7458	0.1037
patch_duplication	0.5001	0.0048
stripes	Not defined	0.0597

14. Rare-Anomaly Sensitivity and Related Work

Historical method: Expected precision at 5% anomalies: 0.4653. False alerts per 1,000: 37.1.

Related work: Genilotti et al., VAD4Space: Visual Anomaly Detection for Planetary Surface Imagery, 2026. Its feature-based benchmarks are complementary references; the central model here remains a normal-only generator.

15. Requirements, Limitations and Reproducibility

- Normal-only deep generative model: Trained diffusion checkpoint and normal manifests preserved.
- Flag and explain anomalies: Saved heatmaps, scoring thresholds, and explanation panels.
- Documented notebook: notebooks/TCMD_SS_NSSC.ipynb executes full end-to-end evaluation.
- At least three model versions: Discussed in Model Evolution section (V1, V2, V3-Final).

16. References

- 1. HiRISE Anomaly Detection problem statement, pp. 1-3.
- 2. NSSC 2026, IIT Kharagpur - Data Analytics event page.
- 3. Lesnikowski et al. (2024). Automated Discovery of Anomalous Features in Planetary Remote Sensing. IEEE JSTARS.
- 4. Gong et al. (2019). Memorizing Normality to Detect Anomaly. ICCV.
- 5. He et al. (2022). Masked Autoencoders Are Scalable Vision Learners. CVPR.
- 6. Wyatt et al. (2022). AnoDDPM. CVPR Workshops.
- 7. Beizaee et al. (2025). DeCo-Diff: Correcting Deviations from Normality. CVPR.
- 8. Guo et al. (2025). Dinomaly. CVPR.
- 9. Zuo et al. (2026). ContaminationAD. Neural Networks.
- 10. Gong et al. (2025). FE-CLIP. ICCV.
- 11. USGS Astrogeology - HiRISE Level 1 processing documentation.
- 12. Meta AI Research - DINOv3.
- 13. Li et al. (2021). CutPaste. CVPR.
- 14. NASA/JPL / Zenodo - HiRISE labeled data set v3.2.